

```

1 import torch, torchvision
2 from torch import nn
3 from torch import optim
4 import torchvision.transforms as TT
5 import torch.nn.functional as F
6 import matplotlib.pyplot as plt

1 from sklearn.metrics import confusion_matrix
2 import pandas as pd
3 import numpy as np
4 from PIL import Image

1 T = torchvision.transforms.Compose([
2     TT.ToTensor(), TT.Normalize((0.5,),(0.5,))
3 ])
4 no_of_batch = 64
5 train_data = torchvision.datasets.MNIST('mnist_data', train=True, download=True)
6 val_data = torchvision.datasets.MNIST('mnist_data', train=False, download=True,
7 train_dl = torch.utils.data.DataLoader(train_data, batch_size = no_of_batch)
8 val_dl = torch.utils.data.DataLoader(val_data, batch_size = no_of_batch)

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
 100% 9912422/9912422 [00:00<00:00,

130235441.74it/s]

Extracting mnist_data/MNIST/raw/train-images-idx3-ubyte.gz to mnist_data/MN

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
 100% 28881/28881 [00:00<00:00, 549959.34it/s]

Extracting mnist_data/MNIST/raw/train-labels-idx1-ubyte.gz to mnist_data/MN

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to
 100% 1648877/1648877 [00:00<00:00,

45074928.77it/s]

Extracting mnist_data/MNIST/raw/t10k-images-idx3-ubyte.gz to mnist_data/MNI

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to



```
1 print(type(train_data[0][0]))
```

```
<class 'torch.Tensor'>
```

```
1 len(train_data)
```

```
60000
```

```
1 len(val_data)
```

```
10000
```

```
1 # for i in range(10):
2 #   t = train_data[i][0]    # t is object containing both image and its label
3 #   img = np.array(t)
4 #   im = plt.imshow(img)
5 #   plt.show()
```

```
1 # Creating a model
2 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
3 class Net(nn.Module):
4     """order of build model:- ConvNet -> Max_Pool -> RELU -> ConvNet -> Max_Pool"""
5     def __init__(self):
6         super(Net, self).__init__()
7         self.conv1 = nn.Conv2d(1, 20, 5, 1)
8         self.conv2 = nn.Conv2d(20, 50, 5, 1)
9         self.fc1 = nn.Linear(4*4*50, 500)
10        self.fc2 = nn.Linear(500, 10)
11
12    def forward(self, x):
13        x = F.relu(self.conv1(x))
14        x = F.max_pool2d(x, 2, 2)
15        x = F.relu(self.conv2(x))
16        x = F.max_pool2d(x, 2, 2)
17        x = x.view(-1, 4*4*50)
18        x = F.relu(self.fc1(x))
19        x = self.fc2(x)
20        return F.log_softmax(x, dim=1)
```

```
1 n_epochs = 4
2 batch_size_train = 64
3 batch_size_test = 1000
4 learning_rate = 0.01
5 # momentum = 0.5
6 log_interval = 10
```

```
1 model = Net().to(device)
2 optimizer = optim.Adam(model.parameters(), lr=learning_rate)
```

```
1 criterion = nn.CrossEntropyLoss()
2 if torch.cuda.is_available():
3     model = model.cuda()
4     criterion = criterion.cuda()
```

```
1 print(model)
```

```
Net(
```

```

(conv1): Conv2d(1, 20, kernel_size=(5, 5), stride=(1, 1))
(conv2): Conv2d(20, 50, kernel_size=(5, 5), stride=(1, 1))
(fc1): Linear(in_features=800, out_features=500, bias=True)
(fc2): Linear(in_features=500, out_features=10, bias=True)
)

```

```

1 # Model Summary
2 from torchsummary import summary
3 summary(model, input_size=(1,28,28))

```

```

-----
              Layer (type)                Output Shape             Param #
=====
              Conv2d-1                [-1, 20, 24, 24]             520
              Conv2d-2                [-1, 50, 8, 8]            25,050
              Linear-3                  [-1, 500]                400,500
              Linear-4                  [-1, 10]                  5,010
=====
Total params: 431,080
Trainable params: 431,080
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.12
Params size (MB): 1.64
Estimated Total Size (MB): 1.76
-----

```

```

1 train_losses = []
2 train_counter = []
3 val_loss1 = []
4 val_losses = []
5 val_counter = []

1 def train(epoch):
2     model.train()
3     for batch_idx, (data, target) in enumerate(train_dl):
4         data = data.cuda()
5         target = target.cuda()
6         optimizer.zero_grad()
7         output = model(data)
8         loss = criterion(output, target)
9         loss.backward()
10        optimizer.step()
11        if batch_idx % log_interval == 0:
12            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
13                epoch, batch_idx * len(data), len(train_dl.dataset),
14                100. * batch_idx / len(train_dl), loss.item()))
15            train_losses.append(loss.item())
16            train_counter.append((batch_idx*64) + ((epoch-1)*len(train_dl.dataset)))

1 def val_loss(epoch):
2     model.eval()
3     for batch_idx, (data, target) in enumerate(val_dl):

```

```

4     data = data.cuda()
5     target = target.cuda()
6     optimizer.zero_grad()
7     output = model(data)
8     loss = criterion(output, target)
9     loss.backward()
10    optimizer.step()
11    if batch_idx % log_interval == 0:
12        print('Val Epoch: {} [{} / {}] ({:.0f}%) \t Loss: {:.6f}'.format(
13            epoch, batch_idx * len(data), len(val_dl.dataset),
14            100. * batch_idx / len(val_dl), loss.item()))
15        val_loss1.append(loss.item())
16        val_counter.append((batch_idx*64) + ((epoch-1)*len(val_dl.dataset)))

```

```

1 def val_accu():
2     model.eval()
3     val_loss = 0
4     correct = 0
5     with torch.no_grad():
6         for data, target in val_dl:
7             data = data.cuda()
8             target = target.cuda()
9             output = model(data)
10            val_loss += criterion(output, target).item()
11            pred = output.data.max(1, keepdim=True)[1]
12            correct += pred.eq(target.data.view_as(pred)).sum()
13    val_loss /= len(val_dl.dataset)
14    val_losses.append(val_loss)
15    print('\nTest set: Avg. loss: {:.4f}, Accuracy: {} / {} ({:.0f}%) \n'.format(
16        val_loss, correct, len(val_dl.dataset),
17        100. * correct / len(val_dl.dataset)))

```

```

1 for epoch in range(1, n_epochs + 1):
2     train(epoch)
3     val_loss(epoch)
4     val_accu()

```

```

Val Epoch: 3 [2800/10000 (28%)] Loss: 0.099846
Val Epoch: 3 [3200/10000 (32%)] Loss: 0.071830
Val Epoch: 3 [3840/10000 (38%)] Loss: 0.157967
Val Epoch: 3 [4480/10000 (45%)] Loss: 0.075607
Val Epoch: 3 [5120/10000 (51%)] Loss: 0.038551
Val Epoch: 3 [5760/10000 (57%)] Loss: 0.028009
Val Epoch: 3 [6400/10000 (64%)] Loss: 0.011252
Val Epoch: 3 [7040/10000 (70%)] Loss: 0.004285
Val Epoch: 3 [7680/10000 (76%)] Loss: 0.024648
Val Epoch: 3 [8320/10000 (83%)] Loss: 0.127178
Val Epoch: 3 [8960/10000 (89%)] Loss: 0.199457
Val Epoch: 3 [9600/10000 (96%)] Loss: 0.199457

```

Test set: Avg. loss: 0.0013, Accuracy: 9807/10000 (98%)

```

Train Epoch: 4 [0/60000 (0%)] Loss: 0.099846
Train Epoch: 4 [640/60000 (1%)] Loss: 0.020276
Train Epoch: 4 [1280/60000 (2%)] Loss: 0.178354
Train Epoch: 4 [1920/60000 (3%)] Loss: 0.175131

```


Train Epoch: 4	[2560/60000 (4%)]	Loss: 0.126930
Train Epoch: 4	[3200/60000 (5%)]	Loss: 0.038620
Train Epoch: 4	[3840/60000 (6%)]	Loss: 0.091242
Train Epoch: 4	[4480/60000 (7%)]	Loss: 0.099996
Train Epoch: 4	[5120/60000 (9%)]	Loss: 0.109382
Train Epoch: 4	[5760/60000 (10%)]	Loss: 0.109759
Train Epoch: 4	[6400/60000 (11%)]	Loss: 0.047011
Train Epoch: 4	[7040/60000 (12%)]	Loss: 0.290384
Train Epoch: 4	[7680/60000 (13%)]	Loss: 0.118871
Train Epoch: 4	[8320/60000 (14%)]	Loss: 0.000545
Train Epoch: 4	[8960/60000 (15%)]	Loss: 0.258911
Train Epoch: 4	[9600/60000 (16%)]	Loss: 0.135672
Train Epoch: 4	[10240/60000 (17%)]	Loss: 0.251775
Train Epoch: 4	[10880/60000 (18%)]	Loss: 0.143522
Train Epoch: 4	[11520/60000 (19%)]	Loss: 0.157935
Train Epoch: 4	[12160/60000 (20%)]	Loss: 0.024100
Train Epoch: 4	[12800/60000 (21%)]	Loss: 0.072212
Train Epoch: 4	[13440/60000 (22%)]	Loss: 0.038258
Train Epoch: 4	[14080/60000 (23%)]	Loss: 0.018026
Train Epoch: 4	[14720/60000 (25%)]	Loss: 0.151196
Train Epoch: 4	[15360/60000 (26%)]	Loss: 0.010569
Train Epoch: 4	[16000/60000 (27%)]	Loss: 0.180696
Train Epoch: 4	[16640/60000 (28%)]	Loss: 0.145909
Train Epoch: 4	[17280/60000 (29%)]	Loss: 0.000164
Train Epoch: 4	[17920/60000 (30%)]	Loss: 0.171249
Train Epoch: 4	[18560/60000 (31%)]	Loss: 0.141871
Train Epoch: 4	[19200/60000 (32%)]	Loss: 0.282870
Train Epoch: 4	[19840/60000 (33%)]	Loss: 0.182656
Train Epoch: 4	[20480/60000 (34%)]	Loss: 0.029370
Train Epoch: 4	[21120/60000 (35%)]	Loss: 0.087558
Train Epoch: 4	[21760/60000 (36%)]	Loss: 0.033027
Train Epoch: 4	[22400/60000 (37%)]	Loss: 0.028485
Train Epoch: 4	[23040/60000 (38%)]	Loss: 0.121118
Train Epoch: 4	[23680/60000 (39%)]	Loss: 0.188900
Train Epoch: 4	[24320/60000 (41%)]	Loss: 0.009253
Train Epoch: 4	[24960/60000 (42%)]	Loss: 0.004044
Train Epoch: 4	[25600/60000 (43%)]	Loss: 0.174736
Train Epoch: 4	[26240/60000 (44%)]	Loss: 0.117377
Train Epoch: 4	[26880/60000 (45%)]	Loss: 0.141804
Train Epoch: 4	[27520/60000 (46%)]	Loss: 0.297068

```

1 fig = plt.figure()
2 plt.plot(train_counter, train_losses, color='blue')
3 plt.legend(['Train Loss'], loc='upper right')
4 plt.xlabel('number of training examples seen')
5 plt.ylabel('Cross Entropy loss')
6
7 fig = plt.figure()
8 plt.plot(val_counter, val_loss1, color='red')
9 plt.legend(['Val Loss'], loc='upper right')
10 plt.xlabel('number of training examples seen')
11 plt.ylabel('Cross Entropy loss')

```

Text(0, 0.5, 'Cross Entropy loss')



```

1 def predict_dl(model, data):
2     y_pred = []
3     y_true = []
4     for i, (images, labels) in enumerate(data):
5         images = images.cuda()
6         x = model(images)
7         value, pred = torch.max(x, 1)
8         pred = pred.data.cpu()
9         y_pred.extend(list(pred.numpy()))
10        y_true.extend(list(labels.numpy()))
11    return np.array(y_pred), np.array(y_true)

```

```

1 y_pred, y_true = predict_dl(model, val_dl)

```

```

1 pd.DataFrame(confusion_matrix(y_true, y_pred, labels=np.arange(0,10)))

```

	0	1	2	3	4	5	6	7	8	9
0	980	0	0	0	0	0	0	0	0	0
1	0	1130	2	1	0	0	0	2	0	0
2	6	2	1013	4	0	0	4	3	0	0
3	0	0	3	1005	0	1	0	1	0	0
4	1	0	0	0	976	0	1	0	0	4
5	3	0	0	13	0	875	0	0	0	1
6	13	3	0	0	3	1	936	0	2	0
7	3	3	4	2	2	1	0	1003	1	9
8	12	3	10	4	2	4	0	4	922	13
9	11	2	0	3	10	0	0	1	0	982

